

Thread Design Considerations

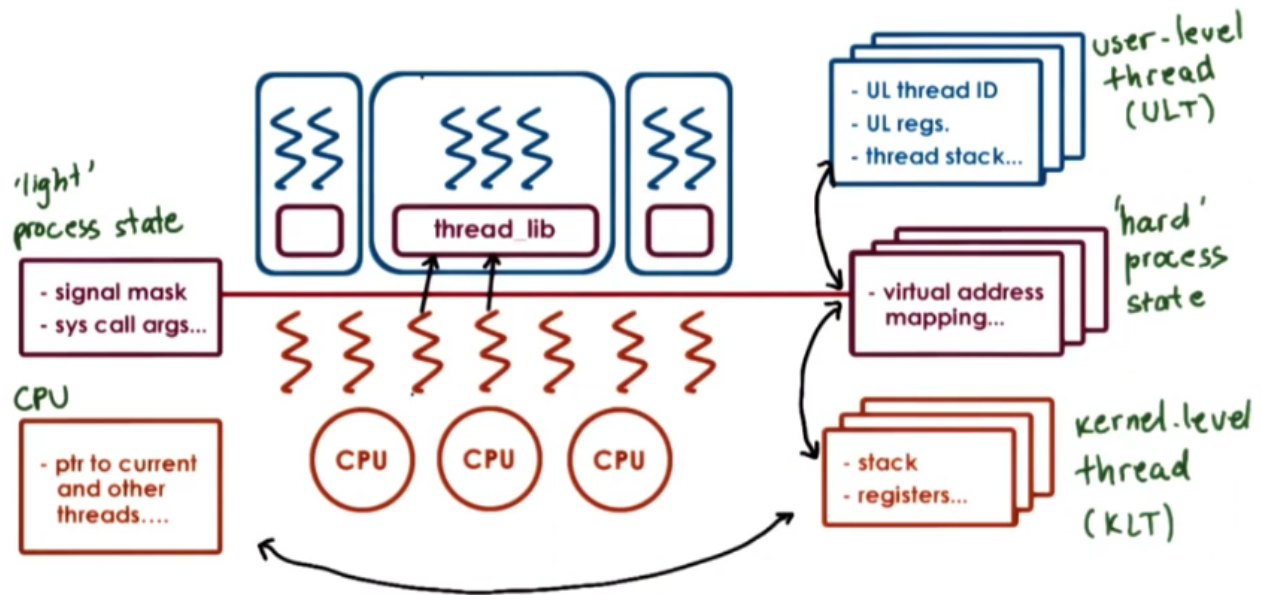
Kernel vs. User Level Threads

1. Both kernel- and user-level threads must provide:
 - Thread abstraction (data structure to define a thread)
 - Scheduling
 - Synchronization
2. The difference is whether these are implemented by the kernel or a library
3. The user-level threads must be mapped to kernel-level threads
 - Must support one-to-one, many-to-one, and many-to-many mappings

Thread Data Structures (Single CPU)

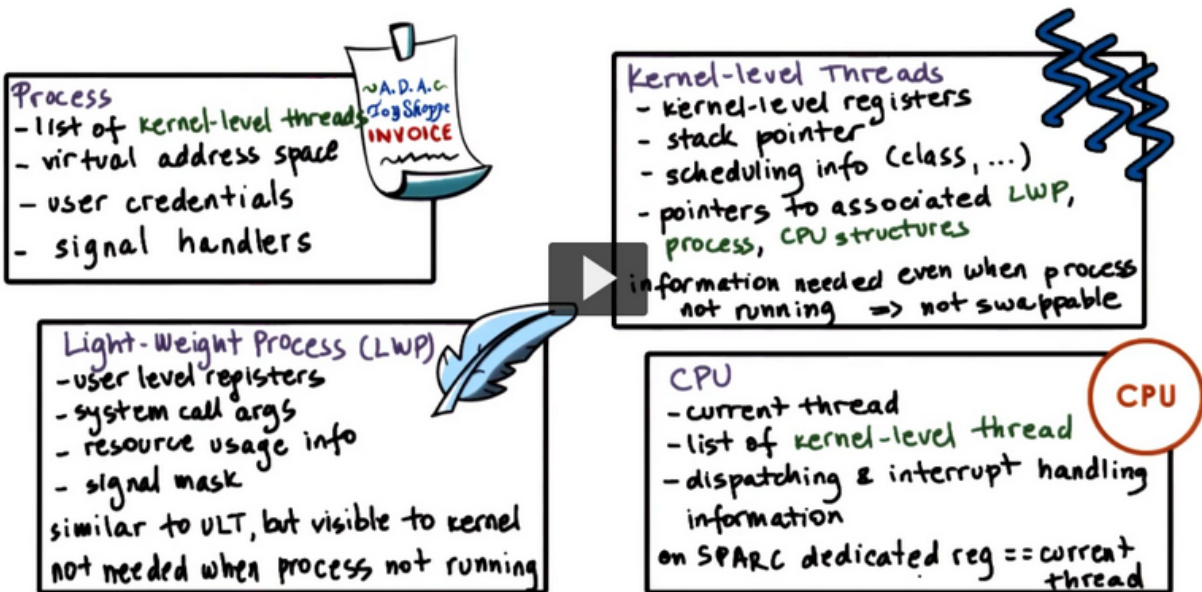
1. Process Control Block (split between user- and kernel-level threads)
 - Virtual address mapping
 - Stack
 - Registers
2. User-Level Thread
 - Thread ID
 - User-level registers
 - Thread stack
3. Kernel-level Thread
 - Stack
 - Registers
4. Light Process State
 - Signal Mask
 - System call arguments
5. Hard Process State
 - Virtual Address Mapping

Thread-related Data structures



Thread-Related Data Structures

Kernel-level Data Structures



Kernel-level Data Structures

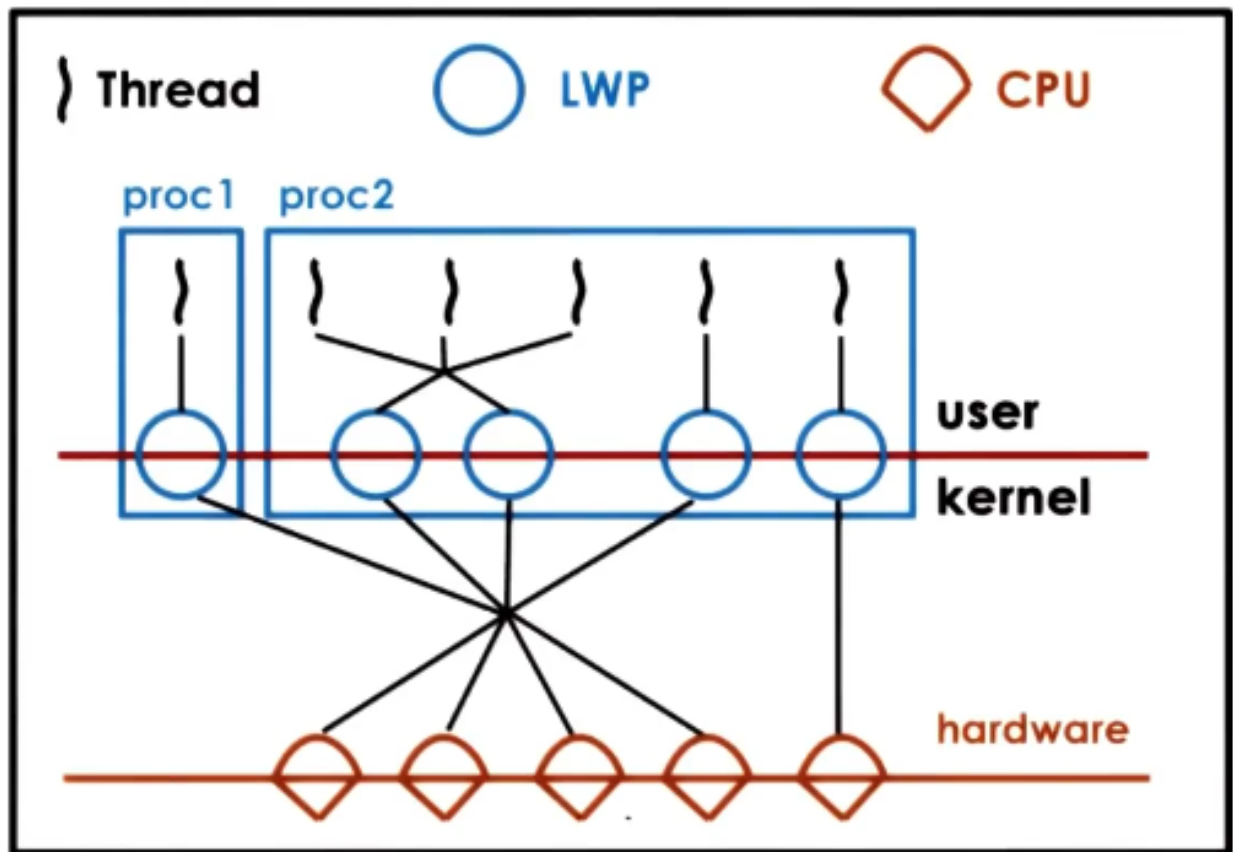
Rationale for Multiple Datastructures

1. Single PCB

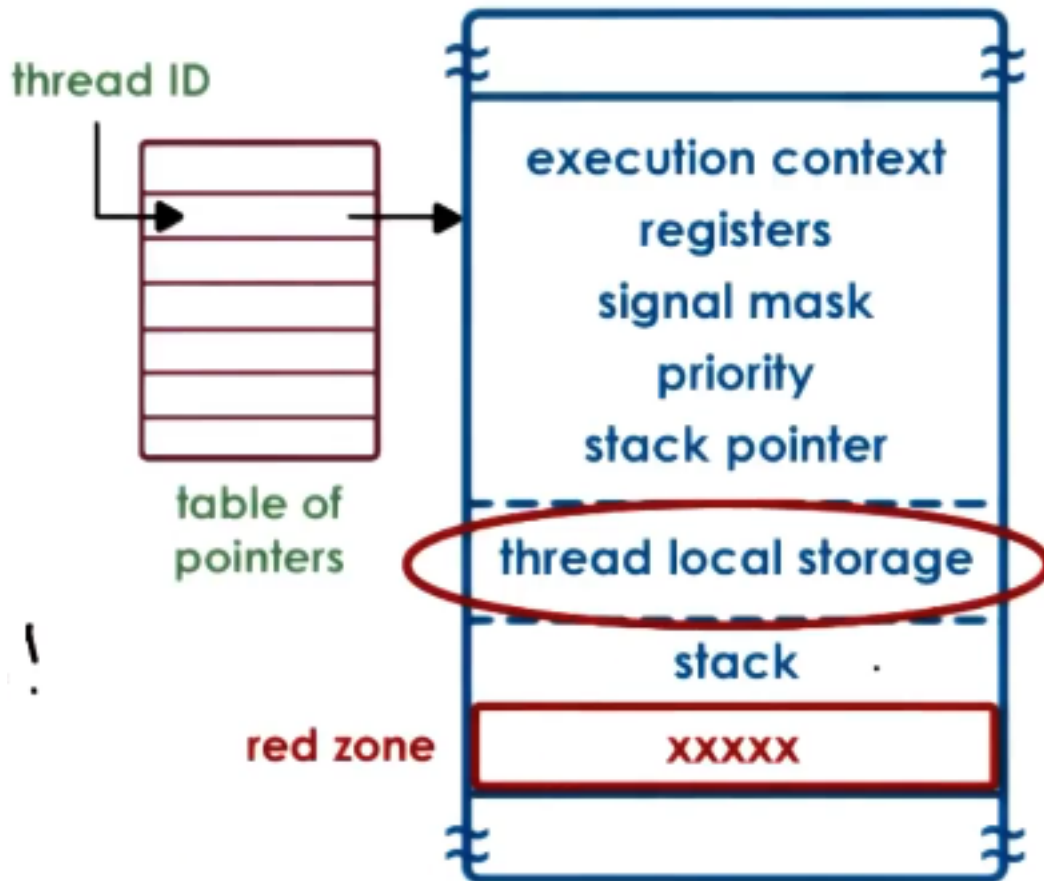
- Large, continuous data structure (limits scalability)
 - Private for each entity (limits overhead)
 - Saved and restored on each context switch (limits performance)
 - Update for any changes (limits flexibility)
2. Multiple Data Structures
 - Smaller data structures
 - Easier to share (pass a pointer)
 - On context switch, only need to save and restore what needs to change
 - User-level library need only update portion of the state
 - Solves the scalability, overhead, performance, and flexibility issues

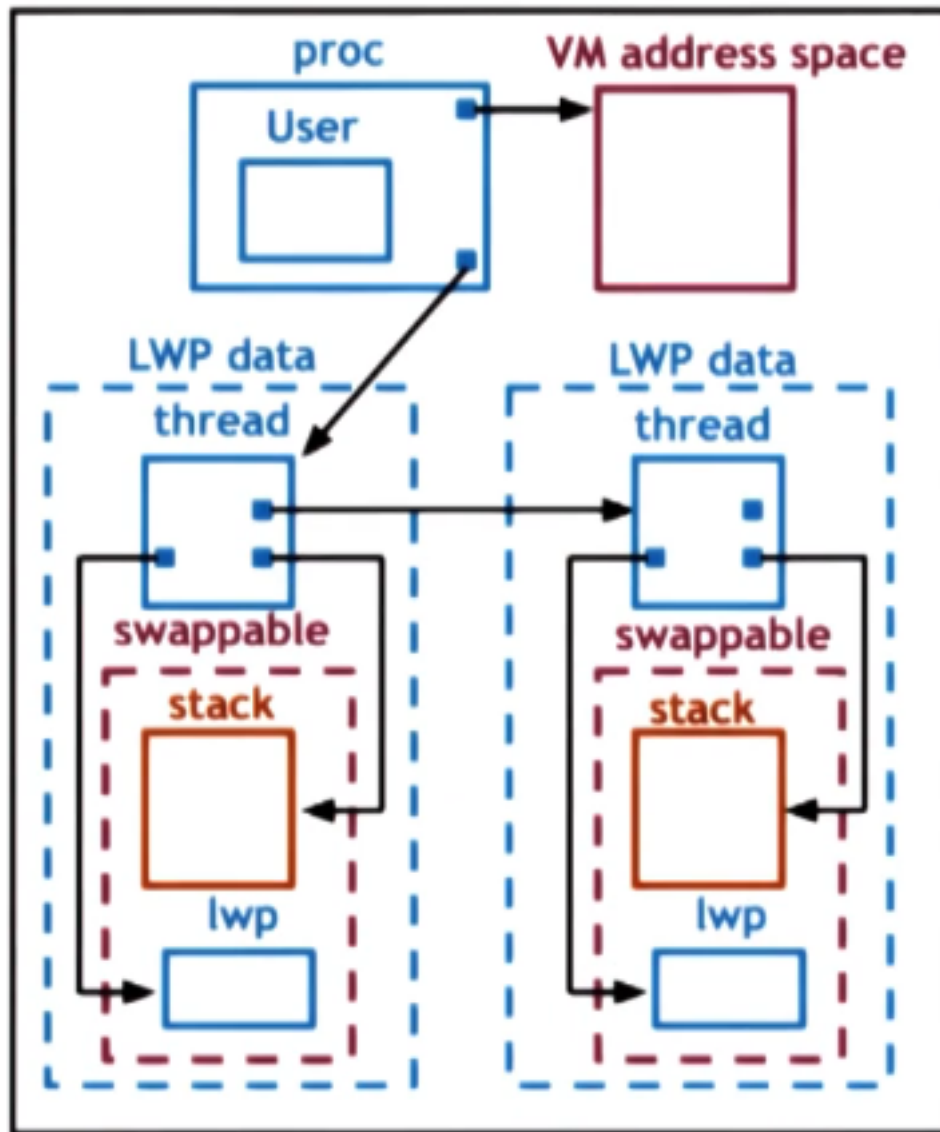
Solaris 2.0 Threading Model

1. “Beyond Multiprocessing: Multithreading the Sun OS Kernel” by Eykholt
2. “Implementing Lightweight Threads” by Stein and Shah
 - Similar to POSIX threads
 - Thread creation -> Thread ID (tid)
 - Table of pointers to thread data structures
 - Stack growth can be dangerous
 - SunOS implemented a “red zone” between blocks of thread memory
3. Kernel-Level Data Structures
 - Process
 - List of kernel-level threads
 - Virtual Address Space
 - User credentials
 - Signal handlers
 - Light-weight Process (LWP)
 - User-level registers
 - System call arguments
 - Resource usage information
 - Signal mask
 - Similar to user-level thread, but visible to kernel
 - Not needed when process not running
 - Kernel-level Threads
 - Kernel-level registers
 - Stack pointer
 - Scheduling information (class, CPU, ...)
 - Pointers to associated LWP, process, CPU structures
 - Information needed even when process isn’t running (not swappable)
 - CPU
 - Current thread
 - List of other kernel-level threads
 - Dispatching and interrupt handling information



Solaris 2.0 Threading Model





Solaris 2.0 Kernel-Level Data Structures

Basic Thread Management Interactions

1. Kernel doesn't know what's happening in user-level library
2. User-level library doesn't know what's happening in the kernel
3. Need some way to communicate between (i.e., LWP's are blocking)
4. System calls and special signals allow kernel/ULT library to coordinate

Thread Management Visibility and Design

1. User-level Library sees:
 - User-level threads
 - Available kernel-level threads
2. Kernel sees:
 - Kernel-level threads

- CPUs
- Kernel-level scheduler
- 3. Can “bind” a user-level thread to a kernel-level thread
- 4. Process jumps to user-level scheduler when:
 - User-level threads explicitly yield
 - Timer set by user-level library expires
 - User-level threads call library functions like lock/unlock
 - Blocked threads become runnable
 - Signals from timer or kernel
- 5. Scheduling becomes more complex between multiple ULTs and KLTs on many CPUs
- 6. For short critical sections, it’s faster for a thread to spin than block (adaptive mutex)
- 7. Destroying threads
 - Faster to reuse threads than destroy if possible
 - When a thread exits...
 - Put on “death row”
 - Periodically destroyed by reaper thread
 - Otherwise thread structures/stacks are reused (performance gains)

Interrupts and Signals

1. Interrupt - Events generated externally by components other than CPU
 - IO devices, timers, other CPUs
 - Based on the physical platform
 - Appear asynchronously
2. Signals - Events triggered by CPU and software running on it
 - Determined based on the operating system
 - Appear synchronously or asynchronously
3. Both:
 - Have a unique ID depending on hardware or OS
 - Can be masked and disabled/suspended via corresponding mask
 - Per-CPU interrupt mask, per-process signal mask
 - If enabled, trigger corresponding handler
 - Interrupt handler set for entire system by OS
 - Signal handlers set on per process basis, by process
 - Handled in specific ways
 - Interrupt and signal handlers
 - Can be ignored
 - Interrupt/signal mask
 - Expected or unexpected
 - Appear synchronously or asynchronously

Interrupt Handling

1. MSI - Message/Signal Interrupt sent over connection (PCI-E)
2. Interrupts are defined by hardware, handling is defined by OS
3. When an interrupt occurs, if it’s enabled, look up the handler in a table
 - Then, jump to the interrupt service routine (ISR)

Signal Handling

1. Signals are defined by OS, handling is defined by process
2. Handlers/Actions - Terminate, ignore, terminate and core dump, stop, continue
3. Process installs handler - signal(), sigaction()
 - For most signals, some cannot be “caught”

4. Synchronous signals
 - SIGSEGV (access to protected memory)
 - SIGFPE (divide by zero)
 - SIGKILL (kill, id) - Can be directed to a specific thread
5. Asynchronous signals
 - SIGKILL (kill)
 - SIGALARM

Disabling Interrupts and Signals

1. Don't want to get interrupted during critical section (mutex held)
 - Can cause deadlock if handler needs same mutex
2. Solutions:
 - Keep handler code simple (no mutexes)
 - Too restrictive
 - Control interruptions by handler code
 - Use interrupt/signal masks
 - 0 -> disabled, 1 -> enabled
 - If disabled, interrupt/signal will wait for it to enable
3. Interrupt masks are maintained on a per-CPU basis
 - If mask disables interrupt, hardware routing mechanism will not deliver it
4. Signal masks are maintained on a per-execution context (ULT on top of KLT)
 - If mask disables signal, kernel sees and will not interrupt thread
5. On multicore systems, can set interrupt on a per-core basis
 - Avoids overheads and perturbations on all other cores
6. Types of Signals
 - One-shot signals - "n signals pending == 1 signal pending" (at least once)
 - Must be explicitly re-enabled
 - Real Time signals - "n signals raised == n handler calls"
7. Signal Examples
 - SIGINT - Terminal interrupt signal
 - SIGURG - High bandwidth data is available on a socket
 - SIGTTOU - Background process attempting write
 - SIGXFSZ - File size limit exceeded

Interrupts as Threads

1. Can handle interrupts in a separate thread, but dynamic thread creation is expensive
2. Instead, make a dynamic decision
 - If handler doesn't lock -> Execute on interrupted thread's stack
 - If handler can block -> Turn into real thread
3. Optimization - Precreate and preinitialize thread structures for interrupt routines
4. Top half - Fast, non-blocking, minimum amount of processing (when interrupt occurs)
5. Bottom half - Arbitrary complexity (can be scheduled for a later time)
6. This is all motivated by performance
 - Overall cost - overhead of 40 SPARC instructions per interrupt
 - Saving of 12 instructions per mutex
 - No changes in interrupt mask, level, ...
 - Fewer interrupts than mutex lock/unlock operations
 - Optimize for the common case

Threads and Signal Handling

1. Need to define a policy for a ULT to tell a KLT that a signal is masked

2. If ULT mask == 1 and KLT mask == 1
 - No issue, signal handling occurs
3. If ULT1 mask == 0 and KLT1 mask == 1, ULT2 mask == 1
 - Library has a signal handler wrapper that intercepts all signals
 - Threading library sends signals to appropriate threads when received
 - ULT1 won't receive but ULT2 will
4. If ULT1 mask == 0 and KLT1 mask == 1, ULT2 mask == 1 and KLT2 mask == 1
 - Signal originates in KLT1, thread_lib knows ULT1 ignores but ULT2 doesn't
 - thread_lib generates directed signal to KLT2 for ULT2
 - KLT2 sends signal to ULT2
5. IF ULT1 mask == 0 and KLT1 mask == 1, ULT2 mask == 0 and KLT2 mask == 1
 - After sending signal, KLT1 will be changed to 0 (thread_lib ignores)
 - The OS will go to KLT2 since it's mask == 1 and try again
 - The process repeats until all KLT mask == 0
6. This handling optimizes the common case
 - Signals are less frequent than signal mask updates
 - System calls avoided - Cheaper to update UL mask
 - Signal handling is more expensive

Tasks in Linux

1. Main execution abstraction is a task
 - Essentially a kernel-level thread
 - Single threaded process has one task
 - Multithreaded process has many tasks
2. Task creation: Clone
 - clone(function, stack_ptr, sharing_flags, arg)
 - fork is implemented as clone with sharing_flags cleared
3. Native POSIX Threads Library (NPTL) "1:1 model"
 - Kernel sees each ULT info
 - Kernel traps are cheaper
 - More resources: memory, large range of IDs
 - Will still break down with very large number of threads (exascale computing)
 - Older LinuxThreads was a many-to-many model
 - Similar issues to those described in Solaris papers

```
struct task_struct {  
    // ...  
    pid_t pid;  
    pid_t tgid;  
    int prio;  
    volatile long state;  
    struct mm_struct *mm;  
    struct files_struct *files;  
    struct list_head tasks;  
    int on_cpu;  
    cpumask_t cpus_allowed;  
    // ...  
}
```

Task Struct in Linux
